

MAD

Decision Making

Learning
Complexe Behavior

Guillaume Lozenguez
@imt-nord-europe.fr



IMT Nord Europe
École Mines-Télécom
IMT-Université de Lille

1. **Back to Q-Learning on 421**
2. **Converging**
3. **Risky Game**
4. **Curse of Dimensionality**



1. **Back to Q-Learning on 421**
2. Converging
3. Risky Game
4. Curse of Dimensionality

Hypotesis: it's Markovian

The system to control
matches a Markov Decision Process

MDP: $\langle S, A, T, R \rangle$:


S : set of system's states

A : set of possible actions

T : $S \times A \times S \rightarrow [0, 1]$: transitions

R : $S \times A \rightarrow R$: cost/rewards

We do have S and A
but not t and r



Q-Learning: the basics

- ▶ Iterative update on (**state**, **action**) evaluation
- ▶ Q-Value equation:

$$Q(s^t, a) = (1 - \alpha)Q(s^t, a) + \alpha \left(r + \gamma \max_{a' \in A} Q(s^{t+1}, a') \right)$$

- ▶ Few parameters:
 α learning rate ; ϵ Exploration-Exploitation ratio and γ discount factor.

Q-Learning: for instance

- ▶ Reaching 4-2-1 at $h-1$ from 6-2-1 at $h-2$ by doing *roll-keep-keep*.




$$Q(2-6-2-1, \text{r-k-k}) = (1 - \alpha)Q(2-6-2-1, \text{r-k-k}) + \alpha \left(r + \gamma \max_{a' \in A} Q(1-4-2-1, a') \right)$$

$$Q(2-6-2-1, \text{r-k-k}) = (1 - \alpha) 40.0 + \alpha (0.0 + 80.0) \quad (a' = \text{keep}^3)$$

- ▶ With α learning rate at 0.1, $Q(2-6-2-1, \text{r-k-k})$ is now equals to 44

Q-Learning: the basics

- ▶ With 500 steps of 500 games:



- ▶ $\alpha: 0.1$; $\epsilon: 0.1$; $\gamma: 0.99$;

Drawing plot in Python: pyplot

Codes:

```
import matplotlib.pyplot as plt


..

plt.plot( values )
plt.ylabel( "mean of the y value" )
plt.show()
```

► Where `values` is a list of values in \mathbb{R}

Q-Learning: the basics




- ▶ With 500 steps of 500 games:



- ▶ With optimal threshold

Q-Learning: the basics




- ▶ With 500 steps of 500 games:



- ▶ α : 0.01 ; ϵ : 0.1 ; γ : 0.99 ;

Q-Learning: the basics

- ▶ With 500 steps of 500 games:



- ▶ α : 0.01 ; ϵ : 0.6 ; γ : 0.99 ;

Playing with the parameters:

- ▶ Generate rapidly "good" policies
- ▶ Converge on a maximal and stable Q-Values
(an indicator for optimal policy)
- ▶ Potentially: be reactive to system modification (recovery)

Ideally: implement dynamic parameters

1. Back to Q-Learning on 421
2. **Converging**
3. Risky Game
4. Curse of Dimensionality

1. Back to Q-Learning on 421
2. **Converging**
3. Risky Game
4. Curse of Dimensionality

Converging

Cf. tutorial:

bitbucket.org/imt-mobisyst/hackagames/src/master/doc/index.md

1. Back to Q-Learning on 421
2. Converging
3. **Risky Game**
4. Curse of Dimensionality

1. Back to Q-Learning on 421
2. Converging
3. Risky Game
4. **Curse of Dimensionality**

The Curse of Dimensionality

1. The Curse
2. Geometric Reduction
3. State Decomposition



1. **The Curse of Dimensionality**

— Example With 2 player 421

2. Geometric Reduction

3. State Decomposition



System Difficulty

Directly correlated to the state space:

The number of states: the Cartesian product of variable domains $|S|$
(minus some unreachable states)

▶ **421 game:** 3 dice-6 at the horizon 3: $(3 \times 6^3 = 648)$ but 168 effective.

Then the branching:

Finally, the number of games:

System Difficulty

Directly correlated to the state space

Then the branching:

The number of possible actions and actions' outcomes.

▶ **421 game:** 2^3 actions, 6^r action outcomes (r , the number of rolled dice - *max. 3*).

Finally, the number of games:

The number of all possible succession of states until reaching an end.

$|Branching|^h$ (h the horizon) - Potentially $|S|^h$. **421 game:** $(6^3)^3$ games

Reminder over Combinatorics

With a Classical 32-card game: Possible distribution $32! = 2.6 \times 10^{35}$



Human life: around 5×10^7 seconds

Probability to play 2 times the same distribution in a human life is very close to 0

Learning 2-players-421

State space ?

Branching ?

First results..

Learning in Combinatorics Context

The root problem: handle large systems

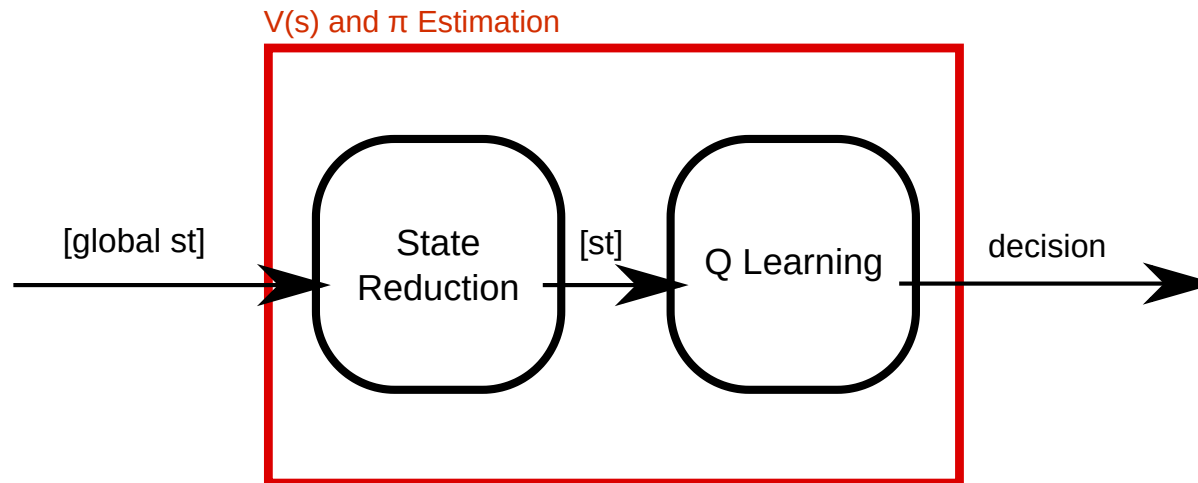
- ▶ Evaluate states $V : S \rightarrow \mathbb{R}$ or $Q : S \times A \rightarrow \mathbb{R}$
- ▶ Build a policy $\pi : S \rightarrow A$

A first basic solution:

- ▶ Reduce the state space definition

State reduction in QLearning

Project the states in a smallest space (dimension and size)




By mitigate the negative impact on the resulting built policy.

1. The Curse of Dimensionality
2. **Geometric Reduction**
 - Reduce the dimension (PCA)
 - Clustering (K-means)
3. State Decomposition

Geometry Reduction

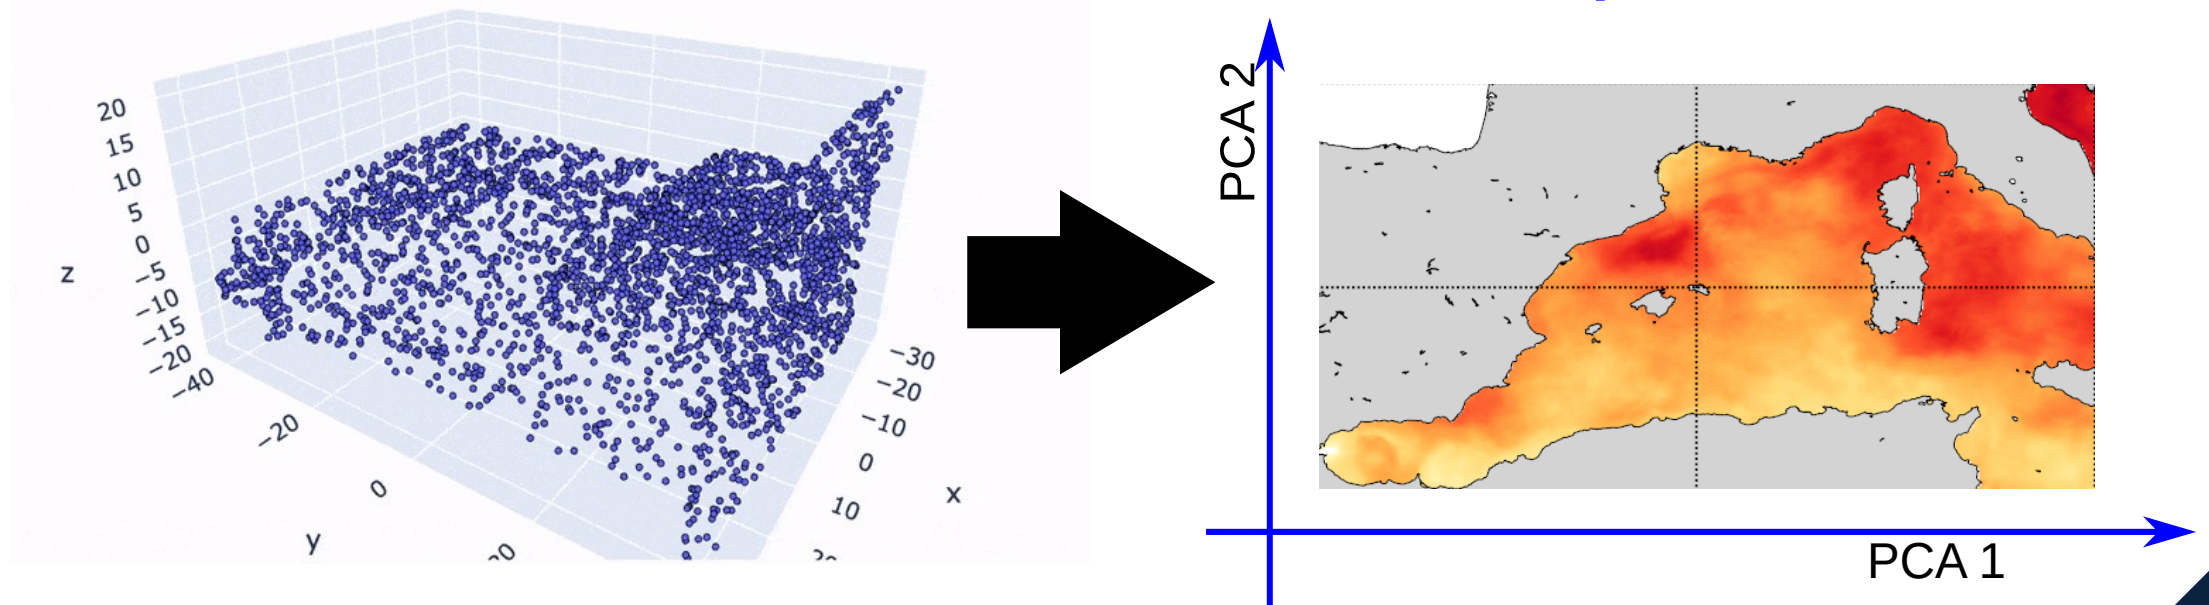
- ▶ Consider that **close** states are similar.
- ▶ Based on the assumption that: *it is possible to define a distance between States*
- ▶ By using regular discretization or adaptative clustering



Reduce the dimension - (Principal Component Analysis)


Searching the hyper-plan that better separate the data, in a given dimension.

State Space



Clustering - (K-means)

regroup the states in coherent sets



K-means:

Searching the optimal k center positions that better group/separate the data

Basic 'simple' classification method

Principal Component Analysis (PCA)

Searching the hyper-plan that better separate the data, in a given dimension.

Python scikit-learn module: **sklearn.decomposition.PCA**

K-means

Searching the optimal k center positions that better group the data together.

Python scikit-learn module: **sklearn.cluster.KMeans**

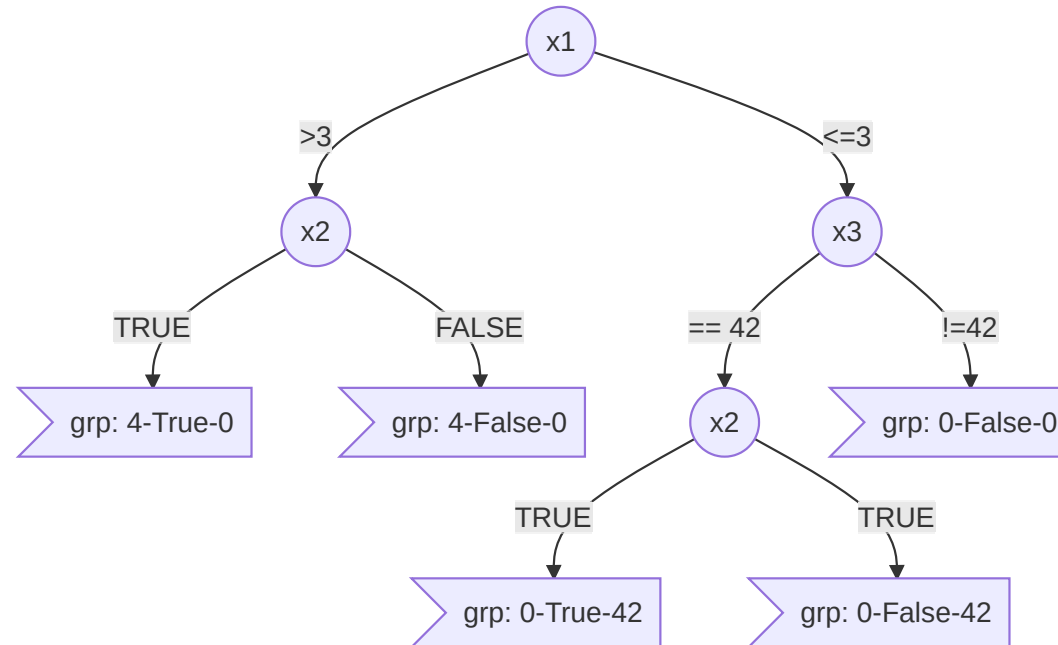
- ▶ Work well with 'linear state transitions' and different states density.
- ▶ Suppose a data set (trace) ideally with proper values

1. The Curse of Dimensionality
2. Geometric Reduction
3. **State-Space Decomposition**
 - Decision Tree
 - Example With 421


State-Space Decomposition

Factorized method: Based on state variable prevalence

- ▶ Decision tree (Again) **Nodes:** variables ; **Edges:** assignment ; **leaf:** group of states



Decision Tree On 421 Q-Learning





Decision Tree On 421 Q-Learning

Simply reduce the state definition to 7 states..

```
def state(self):  
    if self.turn == 0 :  
        return 'end'  
    if self.dices[2] == 1 :  
        if self.dices[1] == 2 :  
            if self.dices[0] == 4 :  
                return "4-2-1"  
            return "X-2-1"  
        if self.dices[1] == 1 :  
            return "X-1-1"  
    return "X-X-1"  
return "X-X-X"
```

Decision Tree On 421 Q-Learning

Results:



Decision Tree Conclusion..

Conclusion:

It is all about defining the appropriate variable prevalence (Decision Tree Structure)

Learn the structure:

- ▶ Expert based Decision Trees or learned ([ID3 algorithm](#))
- ▶ Again on python scikit learn: ([module tree](#))

But..

The evaluation of the structure of the tree is performed by
deadly execution of Q-Learning !