

Q-Learning

A classical method of
Reinforcement Learning

Guillaume Lozenguez

@imt-lille-douai.fr



IMT Nord Europe
École Mines-Télécom
IMT-Université de Lille

Acting over a system evolving under uncertainty

- ▶ **States:** set of configurations defining the studied system
- ▶ **Action:** finite set of possible actions to perform
- ▶ **Transitions:** Describe the possible evolution of the system state

Transition function:

The probabilistic evolution depends on the performed action.

$$T : S \times A \times S \rightarrow [0, 1]$$

$T(s_t, a, s_{t+1})$ return the probability to reach s_{t+1} by doing a from s_t :

$$T(s_t, a, s_{t+1}) = P(s_{t+1} | s_t, a)$$

Acting to optimize Gain

Require to evaluate the interest of each action on the system evolution:

► *Reward/Cost function* (R) :

$$R : S \times A \times S \rightarrow \mathbb{R}$$

$R(s_t, a, s_{t+1})$ is the reward by reaching s_{t+1} from doing a in s_t

OR, in a simplified version:

$$R : S \times A \rightarrow \mathbb{R}$$

Acting to optimize gain (accumulated rewards)

- ▶ Our objective: *a policy* (π) : a function returning the action to perform considering the current state of the system:

$$\pi : S \rightarrow A$$

$\pi(s)$: the action to perform is s

- ▶ *Bellman Equation* :

$$V^\pi(s) = R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') \times V^\pi(s')$$

with : $a = \pi(s)$ and $\gamma \in [0, 1[$ the discount factor (typically 0.99)

reward in 421-game

Over the final combination only with the action "*keep-keep-keep*" or when the horizon is 0

$$\text{score}(4-2-1) = 800$$

$$\text{score}(1-1-1) = 700$$

$$\text{score}(x-1-1) = 400 + x$$

$$\text{score}(x-x-x) = 300 + x$$

$$\text{score}((x+2)-(x+1)-x) = 202 + x$$

$$\text{score}(2-2-1) = 0$$

$$\text{score}(x-x-y) = 100 + x$$

$$\text{score}(y-x-x) = 100 + y$$

Markov Decision Process

MDP: $\langle S, A, T, R \rangle$:

S : set of system's states


A : set of possible actions

T : $S \times A \times S \rightarrow [0, 1]$: transitions

R : $S \times A \rightarrow \mathbb{R}$: cost/rewards

Optimal policy:

The policy π^* maximizing Bellman



Reinforcement Learning:

Learn the optimal policy

- ▶ Without knowledge over the transition probabilities and/or the rewards,
- ▶ but, by getting feedback from acting randomly.

2 approaches

- ▶ **model-based:** Learn the model, then compute the optimal policy.
- ▶ **model-free:** Learn the policy directly.

Model-Free Approaches

Concept

- ▶ Learn without generating **transition** and **reward** models.
- ▶ Build the **policy** directly from the interactions
- ▶ Use only the experience of sequences:

state, action, reward, state, action, ...

Common approaches:

- ▶ **Q-learning**: continuous computing of an expected gain (require rich feedback)
- ▶ **Monte-Carlo**: use random explorations until a 'finale' state (slow to converge).

Exploration–Exploitation tradeoff dilemma

The agent build an optimal behavior from trials and errors.

▶ *Exploration*

- Try new actions to learn unknown feedback
- Better understand the dynamics of the system
- Risky output

▶ *Exploitation*

- Use the best-known action
- Potentially suboptimal

Exploration–Exploitation Tradeoff Dilemma

Examples:

- ▶ *Exploitation*: apply a known game strategy vs *Exploration* investigate new actions.
- ▶ *Exploitation*: go to your favorite restaurant vs *Exploration* try a new one.

Classical approach:

- ▶ Trigger exploration *when* the old fashion strategy doesn't work anymore
Problems:
 - Determine that "a strategy doesn't work" ?
 - Determine that "a new policy is well defined" (exploration end) ?
- ▶ Continuously Explore and Exploite with a fixed ratio.
 - (take wrong decision periodically)

Continuous Exploration–Exploitation : ϵ -Greedy

A Simple heuristic for the Exploration–Exploitation Tradeoff Dilemma

- ▶ Random decision with:
 - a probability ϵ to choose a random action (exploration)
 - a probability $1 - \epsilon$ to choose the best-known action (exploitation)
- ▶ Classically ϵ is set to 0.1
- ▶ A ϵ -greedy agent behavior punctually takes off-policy action

Then the challenge consists in varying ϵ depending of the knowledge the agent has of the area he is interacting in.

Q-learning

One of the most important discoveries in Reinforcement Learning (simple and efficient)

- ▶ At each step, **Q-learning** updates the value attached to a couple (state, action)
- ▶ Updates are performed integrate expected future gains
- ▶ The update is performed accordingly to a learning rate $\alpha \in]0, 1[$
→ α : ratio between new vs old accumulated information.

Q-learning based on a Q function

Considering it is not possible to evaluate state without a policy yet

$$V^\pi(s) = R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') \times V^\pi(s')$$

the **Q-values** evaluate each action performed from each state:

$$Q : S \times A \rightarrow \mathbb{R}, \quad Q(s, a) \text{ is the value of doing } a \text{ from } s$$

and, a **Q-value** is updated iteratively from succession of: $\langle s, a, s', r \rangle$

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha \left(r + \gamma \max_{a' \in A} Q(s', a') \right)$$

Q-learning : the algorithm

Input: state and action spaces: A ; a step engine *Perform* ;
exploration ratio: ϵ ; learning rate: α ; discount factor γ

1. Read the initial state s
2. Initialize $Q(s, a)$ to 0 for any action a
3. Repeat until convergence
 1. At ϵ random: get a random a *or* a maximizing $Q(s, a)$
 2. *Perform* a and read the reached state s' and the associated reward r
 3. If necessary, add s' to Q (with value 0 for any action a)
 4. Update $Q(s, a)$ accordingly to α and γ
 5. set $s = s'$

Output: the **Q-values**.

Q-learning : the algorithm

In agent-based programming:

- ▶ As an initial step :
 1. Initialize Q
- ▶ At 'game' start :
 1. Read the initial state s
- ▶ At each iteration :
 1. Read the reached state s' and the associated reward r
 2. If necessary, add s' to Q (with value 0 for any action a)
 3. Update $Q(s, a)$ accordingly to α and γ
 4. record $s = s'$
 5. At ϵ random: get a random a *or* a maximizing $Q(s, a)$

Q-learning : the main equation

Update Q each time a tuple $\langle s^t, a, s^{t+1}, r \rangle$ is read

$$newQ(s^t, a) = (1 - \alpha)Q(s^t, a) + \alpha \left(r + \gamma \max_{a' \in A} Q(s^{t+1}, a') \right)$$

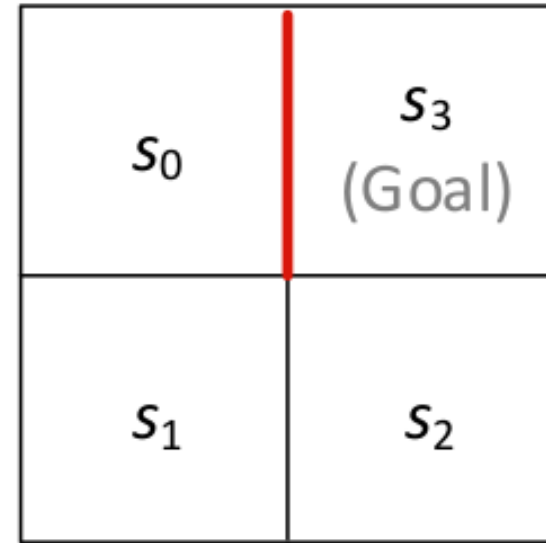
- ▶ $Q : S \times A \rightarrow \mathbb{R}$: the value function we build.
- ▶ α : the learning rate, ϵ : the Exploration-Exploitation ratio
- ▶ γ : the discount factor

The known optimal policy:

$$\pi^*(s) = \max_{a \in A} Q(s, a)$$

Simple Example

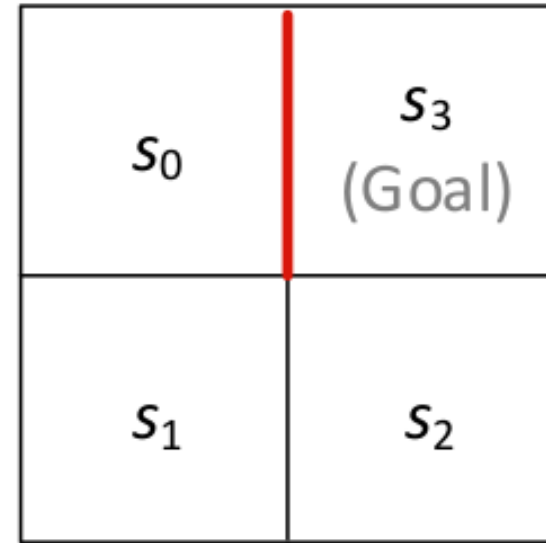
- ▶ **States:** 4 positions
 s_0 , s_1 , s_2 and s_3
- ▶ **Actions:** left, right, up, down
- ▶ **Transitions:** determinist
- ▶ **Rewards:**
10 for reaching s_3 , -1 else



($\epsilon = 0.1$, $\alpha = 0.1$ and $\gamma = 0.99$)

Simple Example

- ▶ From s_0 get action *left* (explore)
reaches s_0 with -1
updates $Q(s_0, \textit{left}) = -0.1$
- ▶ s_0 gets *right* (best) $\rightarrow (s_0, -1)$
updates $Q(s_0, \textit{right}) = -0.1$
- ▶ s_0 gets *down* (exp.) $\rightarrow (s_1, -1)$
updates $Q(s_0, \textit{down}) = -0.1$
...
- ▶ s_2 gets *up* (exp.) $\rightarrow (s_3, 10)$
updates $Q(s_2, \textit{up}) = 1$
End Episode



Simple Example

($\alpha = 0.1$, $\epsilon = 0.1$ and $\gamma = 0.99$)


► **Episode 1:** (18 action)

S	s_0	s_1	s_2
$\max Q$	-0.39	-0.19	1

► **Episode 2:** (15 action)

S	s_0	s_1	s_2
$\max Q$	-0.43	0.9	1.9

...

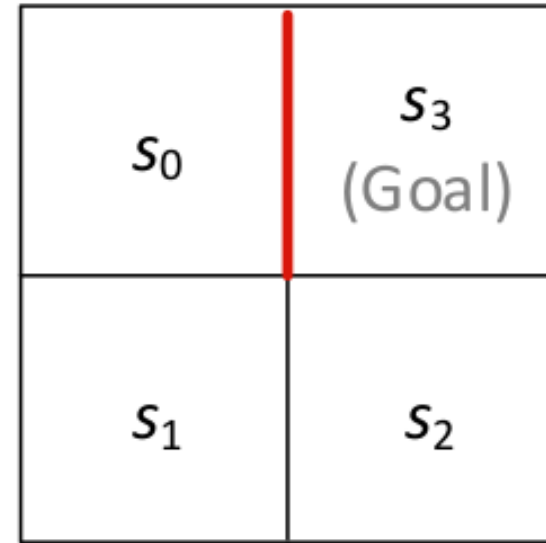


Simple Example

($\alpha = 0.1$, $\epsilon = 0.1$ and $\gamma = 0.99$)

► **Episode N:** (3-4 actions)

S	s_0	s_1	s_2
$\max Q$	7.8	8.9	10
$\operatorname{argmax} Q$	↓	→	↑





Let's go....