

# Q-Learning

A classical method of  
Reinforcement Learning

Guillaume.Lozenquez

[@imt-nord-europe.fr](mailto:@imt-nord-europe.fr)



**IMT Nord Europe**  
École Mines-Télécom  
IMT-Université de Lille

1. **A teoritical framework: Markov Decision Process**
2. **On the go, model free learning**
  - **Compute QValues**
  - **Choose an Action**
3. **Exercise**

# Acting over a system evolving under uncertainty

- ▶ **States:** set of configurations defining the studied system
- ▶ **Action:** finite set of possible actions to perform
- ▶ **Transitions:** Describe the possible evolution of the system state

## Transition function:

The probabilistic evolution depends on the performed action.

$$T : S \times A \times S \rightarrow [0, 1]$$

$T(s^t, a, s^{t+1})$  return the probability to reach  $s^{t+1}$  by doing  $a$  from  $s^t$ :

$$T(s^t, a, s^{t+1}) = P(s^{t+1} | s^t, a)$$

# Transition in 421-game

► For instance, doing *Keep-Kepp-Roll* in **4-2-2 (2)** :

—  $6-4-2 (1) = 1/6$

—  $5-4-2 (1) = 1/6$

—  $4-4-2 (1) = 1/6$       Or  $T(422(2), \text{k-k-r}, 442(1)) = 1/6$

—  $4-3-2 (1) = 1/6$

—  $4-2-2 (1) = 1/6$

—  $4-2-1 (1) = 1/6$

► For instance, doing *Keep-Kepp-Kepp* in **1-1-1 (2)** :       $1-1-1 (0) = 1$

# Acting to optimize Gain

Require to evaluate the interest of each action on the system evolution:

► *Reward/Cost function* (R) :

$$R : S \times A \times S \rightarrow \mathbb{R}$$

$R(s^t, a, s^{t+1})$  is the reward by reaching  $s^{t+1}$  from doing  $a$  in  $s^t$

**OR**, in a simplified version:

$$R : S \times A \rightarrow \mathbb{R}$$

## reward in 421-game

Over the final combination when the horizon reaches 0

$$\text{score}(4-2-1) = 800$$

$$\text{score}(1-1-1) = 700$$

$$\text{score}(x-1-1) = 400 + x$$

$$\text{score}(x-x-x) = 300 + x$$

$$\text{score}((x+2)-(x+1)-x) = 202 + x$$

$$\text{score}(2-2-1) = 0$$

$$\text{score}(x-x-y) = 100 + x$$

$$\text{score}(y-x-x) = 100 + y$$

**Reward function:**  $r(s, a, s') = \text{score}(s')$  if  $h = 0$ ; 0 else

# Acting to optimize gain (accumulated rewards)

- ▶ Our objective: *a policy* ( $\pi$ ) : a function returning the action to perform considering the current state of the system:

$$\pi : S \rightarrow A$$

$\pi(s)$  : the action to perform is  $s$

- ▶ *Bellman Equation* :

$$V^\pi(s) = R(s^t, a) + \gamma \sum_{s^{t+1} \in S} T(s^t, a, s^{t+1}) \times V^\pi(s^{t+1})$$

with :  $a = \pi(s)$  and  $\gamma \in [0, 1[$  the discount factor (typically 0.99)

# Markov Decision Process

**MDP:**  $\langle S, A, T, R \rangle$ :

$S$ : set of system's states

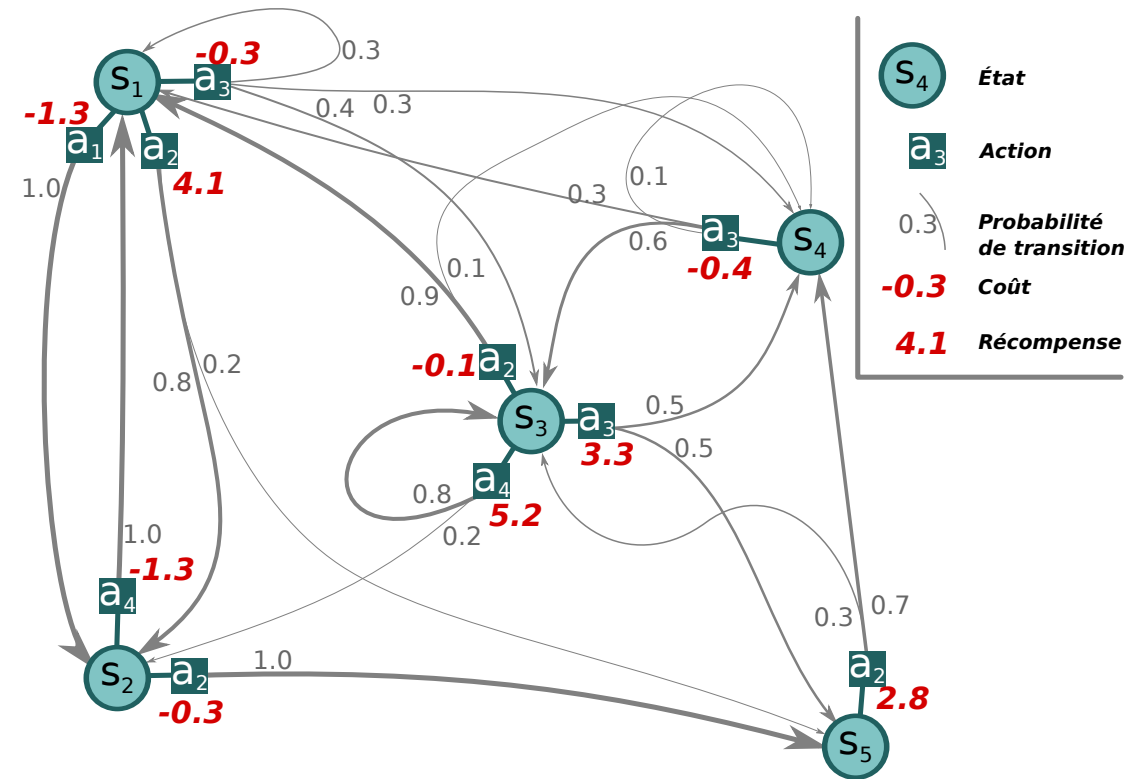
$A$ : set of possible actions

$T: S \times A \times S \rightarrow [0, 1]$ : transitions

$R: S \times A \rightarrow R$ : cost/rewards

**Optimal policy:**

The policy  $\pi^*$  maximizing Bellman





# Reinforcement Learning:

## Learn the optimal policy

- ▶ Without knowledge over the transition probabilities and/or the rewards,
- ▶ but, by getting feedback from acting randomly.

## 2 approaches

- ▶ **model-based:** Learn the model  $(T, R)$ , and compute a policy.
- ▶ **model-free:** Learn the policy directly.

1. A teoritical framework: Markov Decision Process
2. **On the go, model free learning**
  - **Compute QValues**
  - Choose an Action
3. Exercice

# Model-Free Approaches

## Concept

- ▶ Learn without generating **transition** and **reward** models.
- ▶ Build the **policy** directly from the interactions
- ▶ Use only the experience of sequences:  
*state, action, reward, state, action, ...*

## Common approaches:

- ▶ **Q-learning:**  
continuous computing of an expected gain (require rich feedback)
- ▶ **Monte-Carlo:** use random explorations until a 'finale' state (slow to converge).

# Q-learning

One of the core discoveries in Reinforcement Learning (simple and efficient)

- ▶ At each step, **Q-learning** updates the value attached to a couple (state, action)
- ▶ Updates are performed integrate expected future gains
- ▶ The update is performed accordingly to a learning rate  $\alpha \in ]0, 1[$   
→  $\alpha$  : ratio between new vs old accumulated information.

# Q-learning based on a Q function

Considering it is not possible to evaluate state without a policy yet

$$V^{\pi}(s) = R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') \times V^{\pi}(s')$$

the **Q-values** evaluate each action performed from each state:

$$Q : S \times A \rightarrow \mathbb{R}, \quad Q(s, a) \text{ is the value of doing } a \text{ from } s$$

and, a **Q-value** is updated iteratively from succession of:  $\langle s, a, s', r \rangle$

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha (r + \gamma Q(s', a'))$$

# Q-learning : the algorithm

*Input:* state and action spaces:  $A$  ; a step engine *Perform* ;  
exploration ratio:  $\epsilon$  ; learning rate:  $\alpha$  ; discount factor  $\gamma$

1. Read the initial state  $s$
2. Initialize  $Q(s, a)$  to 0 for any action  $a$
3. Repeat until convergence
  - i. Select an action  $a$  with random
  - ii. *Perform*  $a$  and read the reached state  $s'$  and the associated reward  $r$
  - iii. If necessary, add  $s'$  to  $Q$  ( with value 0 for any action  $a$ )
  - iv. Update  $Q(s, a)$  accordingly to  $\alpha$  and  $\gamma$
  - v. Set  $s = s'$

*Output:* the **Q-values**.

# Q-learning : the main equation

Update Q each time a tuple  $\langle s^t, a, s^{t+1}, r \rangle$  is read

$$newQ(s, a) = (1 - \alpha) Q(s, a) + \alpha (\text{incomming-feedback})$$

$$\text{incomming-feedback} = r(s, a, s') + \gamma Q(s', a')$$

- ▶  $\alpha$  : the learning rate (= 0.1)
- ▶  $\gamma$  : the discount factor (= 0.999)

The known optimal policy:

$$\pi^*(s) = \max_{a \in A} Q(s, a)$$

1. A teoritical framework: Markov Decision Process
2. **On the go, model free learning**
  - Compute QValues
  - **Choose an Action**
3. Exercice



# Exploration–Exploitation tradeoff dilemma

The agent build an optimal behavior from trials and errors.

## ▶ *Exploration*

- Try new actions to learn unknown feedback
- Better understand the dynamics of the system
- Risky output

## ▶ *Exploitation*

- Use the best-known action
- Potentially suboptimal

# Exploration–Exploitation Tradeoff Dilemma

## Examples:

- ▶ *Exploitation*: apply a known game strategy **vs** *Exploration* investigate new actions.
- ▶ *Exploitation*: go to your favorite restaurant **vs** *Exploration* try a new one.

## Classical approach:

- ▶ Trigger exploration *when* the old fashion strategy doesn't work anymore
- Problems:
  - Determine that "a strategy doesn't work" ?
  - Determine that "a new policy is well defined" (exploration end) ?
- ▶ Continuously Explore and Exploite with a fixed ratio.
  - (take wrong decision periodically)

# Continuous Exploration–Exploitation : $\epsilon$ -Greedy

A Simple heuristic for the Exploration–Exploitation Tradeoff Dilemma

- ▶ Random decision with:
  - a probability  $\epsilon$  to choose a random action (exploration)
  - a probability  $1 - \epsilon$  to choose the best-known action (exploitation)
- ▶ Classically  $\epsilon$  is set to 0.1
- ▶ A  $\epsilon$ -greedy agent behavior punctually takes off-policy action

Then the challenge consists in varying  $\epsilon$  depending of the knowledge the agent has of the area he is interacting in.

# Q-learning : the algorithm

*Input:* state and action spaces:  $A$  ; a step engine *Perform* ;  
exploration ratio:  $\epsilon$  ; learning rate:  $\alpha$  ; discount factor  $\gamma$

1. Read the initial state  $s$
2. Initialize  $Q(s, a)$  to 0 for any action  $a$
3. Repeat until convergence
  - i. **At  $\epsilon$  random: get a random  $a$  or a maximizing  $Q(s, a)$**
  - ii. *Perform*  $a$  and read the reached state  $s'$  and the associated reward  $r$
  - iii. If necessary, add  $s'$  to  $Q$  ( with value 0 for any action  $a$ )
  - iv. Update  $Q(s, a)$  accordingly to  $\alpha$  and  $\gamma$
  - v. set  $s = s'$

*Output:* the **Q-values**.

# Q-learning : In Agent-Based Architecture

- ▶ As an initial step (**wakeUp**) :
  - Initialize  $Q$
  - Initialize state and action variables ( $s, a$ ).
- ▶ At each iteration (**perceive**):
  - Read the reached state  $s'$  and the associated reward  $r$
  - If necessary, add  $s'$  to  $Q$  (with value 0 for any action  $a$ )
  - Update  $Q(s, a)$  accordingly to  $\alpha$  and  $\gamma$
  - record  $s = s'$
- ▶ Taking decisions (**decide**):
  - At  $\epsilon$  random: get a random  $a$  *or*  $a$  maximizing  $Q(s, a)$

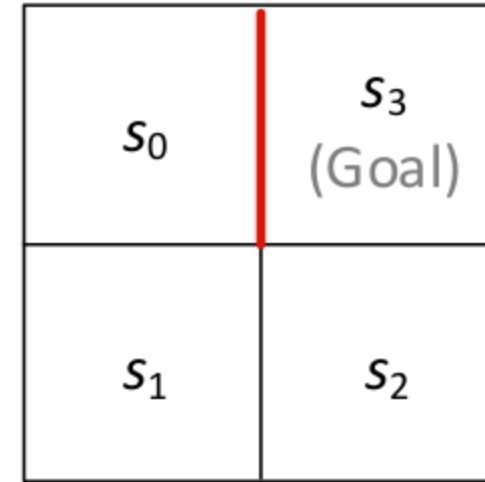
1. A teoritical framework: Markov Decision Process
2. **On the go, model free learning**
  - Compute QValues
  - **Choose an Action**
3. Exercice

Applying Q-Learning...

# Simple Example

- ▶ **States:** 4 positions  
 $s_0, s_1, s_2$  and  $s_3$
- ▶ **Actions:** left, right, up, down
- ▶ **Transitions:** deterministic
- ▶ **Rewards:**  
10 for reaching  $s_3$ , -1 else

( $\epsilon = 0.1, \alpha = 0.1$  and  $\gamma = 0.99$ )

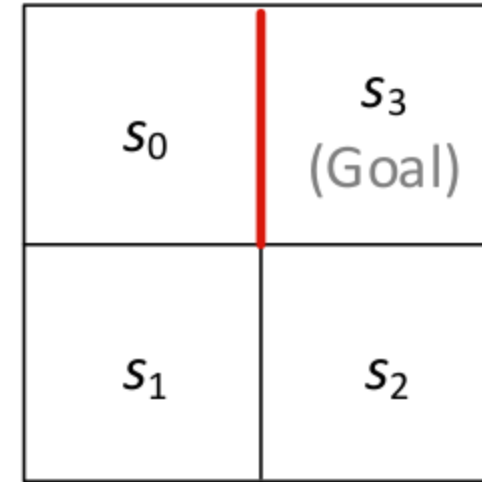


( $\alpha = 0.1, \epsilon = 0.1$  and  $\gamma = 0.99$ )



# Simple Example

- ▶ From  $s_0$  get action *left* (explore)  
reaches  $s_0$  with  $-1$   
updates  $Q(s_0, \textit{left}) = -0.1$
- ▶  $s_0$  gets *right* (best)  $\rightarrow (s_0, -1)$   
updates  $Q(s_0, \textit{right}) = -0.1$
- ▶  $s_0$  gets *down* (exp.)  $\rightarrow (s_1, -1)$   
updates  $Q(s_0, \textit{down}) = -0.1$   
...
- ▶  $s_2$  gets *up* (exp.)  $\rightarrow (s_3, 10)$   
updates  $Q(s_2, \textit{up}) = 1$   
**End Episode**



$(\alpha = 0.1, \epsilon = 0.1 \text{ and } \gamma = 0.99)$

# Simple Example

( $\alpha = 0.1$ ,  $\epsilon = 0.1$  and  $\gamma = 0.99$ )


► **Episode 1:** ( 18 action)

<b>S</b>	$s_0$	$s_1$	$s_2$
$\max Q$	-0.39	-0.19	1

► **Episode 2:** ( 15 action)

<b>S</b>	<b>s_0</b>	<b>s_1</b>	<b>s_2</b>
$\max Q$	-0.43	0.9	1.9

...



( $\alpha = 0.1$ ,  $\epsilon = 0.1$  and  $\gamma = 0.99$ )

# Simple Example

► **Episode N:** (3-4 actions)

<b>S</b>	$s_0$	$s_1$	$s_2$
$\max Q$	7.8	8.9	10
$\operatorname{argmax} Q$	↓	→	↑



( $\alpha = 0.1$ ,  $\epsilon = 0.1$  and  $\gamma = 0.99$ )

# Exercise: Apply Q-Learning

## Agent Version:

- ▶ On 421 game of [hackagame](#)

## Classical version:

- ▶ On Lunar-Lander game of [farama::gymnasium](#)