

Learning 421 game

Model-Based Learning

Guillaume Lozenguez

@imt-nord-europe.fr



IMT Nord Europe
École Mines-Télécom
IMT-Université de Lille

1. **Back to Q-Learning on 421**
2. **ON or OFF policy**
3. **Scale-Up**

Q-Learning: the basics

- ▶ Iterative update on (**state**, **action**) evaluation
- ▶ Q-Value equation:

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha (r + \gamma Q(s', a'))$$

- ▶ Few parameters:
 α learning rate ; ϵ Exploration-Exploitation ratio and γ discount factor.

Q-Learning: for instance

- ▶ Reaching 4-2-1 from 6-2-1 by doing *roll-keep-keep*.




$$Q(6-2-1, \text{r-k-k}) = (1 - \alpha)Q(6-2-1, \text{r-k-k}) + \alpha (r + \gamma Q(4-2-1, a'))$$

$$Q(2-6-2-1, \text{r-k-k}) = (1 - \alpha) 40.0 + \alpha (0.0 + 80.0) \quad (a' = \text{keep}^3)$$

- ▶ With α learning rate at 0.1, $Q(6-2-1, \text{r-k-k})$ is now equals to 44

Q-Learning: the basics

- ▶ With 500 steps of 500 games:



- ▶ $\alpha: 0.1$; $\epsilon: 0.1$; $\gamma: 0.99$;

Drawing plot in Python: pyplot

Codes:

```
import matplotlib.pyplot as plt


...

plt.plot( values )
plt.ylabel( "mean of the y value" )
plt.show()
```

► Where `values` is a list of values in \mathbb{R}

Q-Learning: the basics




- ▶ With 500 steps of 500 games:



- ▶ With optimal threshold

Q-Learning: the basics




- ▶ With 500 steps of 500 games:



- ▶ α : 0.01 ; ϵ : 0.1 ; γ : 0.99 ;

Q-Learning: the basics

- ▶ With 500 steps of 500 games:



- ▶ α : 0.01 ; ϵ : 0.6 ; γ : 0.99 ;

Playing with the parameters:

- ▶ Generate rapidly "good" policies
- ▶ Converge on a maximal and stable Q-Values
(an indicator for optimal policy)
- ▶ Potentially: be reactive to system modification (recovery)

Ideally: implement dynamic parameters

1. Back to Q-Learning on 421
2. **ON or OFF policy**
3. Scale-Up

On or Off Policy: The main idea

Does the Q -values match the actual in-use policy ?

- ▶ **ON** the Q -values and the in-use policy π are aligned.
- ▶ **OFF** the learning Q -values model something else.
 - Generally π is ϵ -greedy and Q -values optimal

On or Off Policy: more formally

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha \left(r + \gamma \max_{a' \in A} Q(s', a') \right)$$

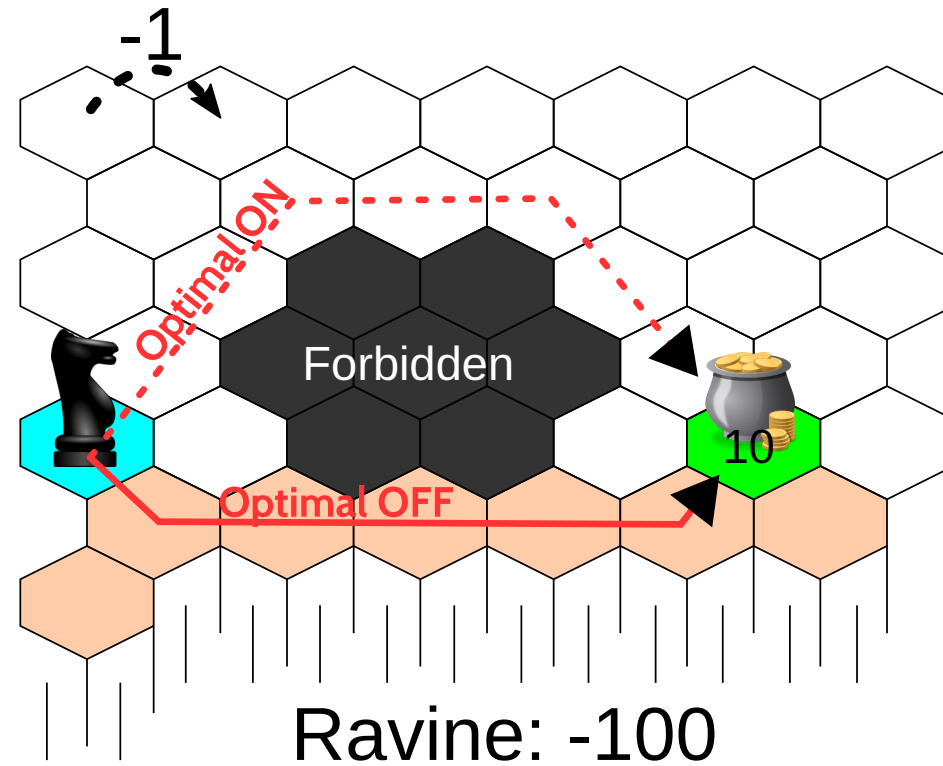
► **ON:**

$a' = \pi(s')$ (with π the ϵ -greedy policy for instance)

► **OFF:**

$$a' = \max_{a' \in A} Q(s', a') \quad \left(\text{and } \pi^*(s) = \arg \max_{a \in A} Q(s, a) \right)$$

On or Off Policy: with an example:



- Deterministic action outcome but 10% to random action.

1. Back to Q-Learning on 421
2. ON or OFF policy
3. **Scale-Up**

Scale-Up:

The central difficulty in Machine-Learning

- ▶ You are encouraged to try *QLearning* on 421 Duo mode...
 - From scratch to win random AI.
 - From 421-Solo to win 421-Solo playing Duo.

Scale-Up:

Heuristic values

- ▶ Expert-oriented reward function.
- ▶ Expert-based qvalue initialization.

Iterative Learning

- ▶ Start with few state and action and grow the model.

Hybrid Approach

- ▶ Process Q-Learning a record in the same time
- ▶ Use punctually heavy 'classical' learning