

AI and Games

Model-Based Learning

Guillaume Lozenguez

@imt-nord-europe.fr



IMT Nord Europe
École Mines-Télécom
IMT-Université de Lille

Model-based learning

Main Idea:

- ▶ Random trajectories (a lot)
- ▶ Until each transition is visited several times.
- ▶ Compute an optimal policy.

Potentially:

- ▶ Require driven exploration to go in every 'niche'
- ▶ But generally: only incomplete exploration can be performed

But first the Model

Markov Decision Process

A framework for modeling stochastic evolution of the system to control.

Bellman equation

Recursive evaluation of states to compute expected gains.

Solving algorithms

- ▶ Value iteration
- ▶ Policy iteration

Markov Decision Process

MDP: $\langle S, A, T, R \rangle$:

S : set of system's states

A : set of possible actions

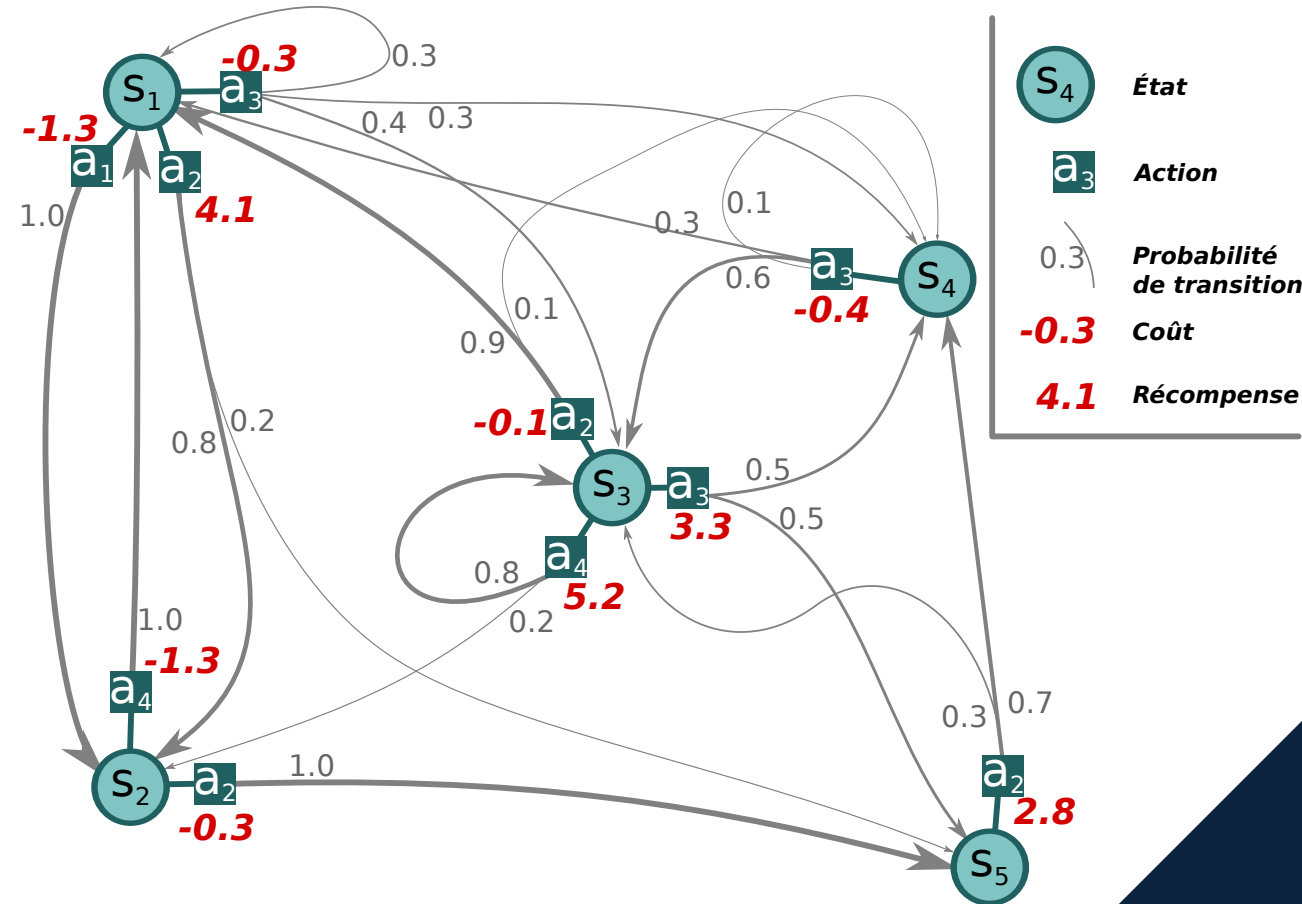
T : $S \times A \times S \rightarrow [0, 1]$: transitions

R : $S \times A \rightarrow \mathbb{R}$: average cost/rewards

Optimal policy:

π : a function returning the action to perform in each crossed states.

π^* : the optimal policy maximizing the gains (expected cumulated rewards).



Bellman Equation

State evaluation for a given policy π :

$$V^\pi(s) = R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') \times V^\pi(s')$$

with : $a = \pi(s)$ and $\gamma \in [0, 1]$ the discount factor (typically 0.99)

As a sum of gains:

- ▶ The immediate reward: $R(s, a)$.
- ▶ The future gains $V^\pi(s')$, for all possible next states $s' \in S$,
- ▶ proportionally to the probability to reach them $T(s, a, s')$

Solving MDP: Value Iteration

Input: an **MDP**: $\langle S, A, T, R \rangle$; precision error: ϵ ; discount factor: γ ; initial $V(s)$

1. Repeat until: **maximal delta** $< \epsilon$

For each state $s \in S$

 — Search the action a^* maximizing the Bellman Equation on s

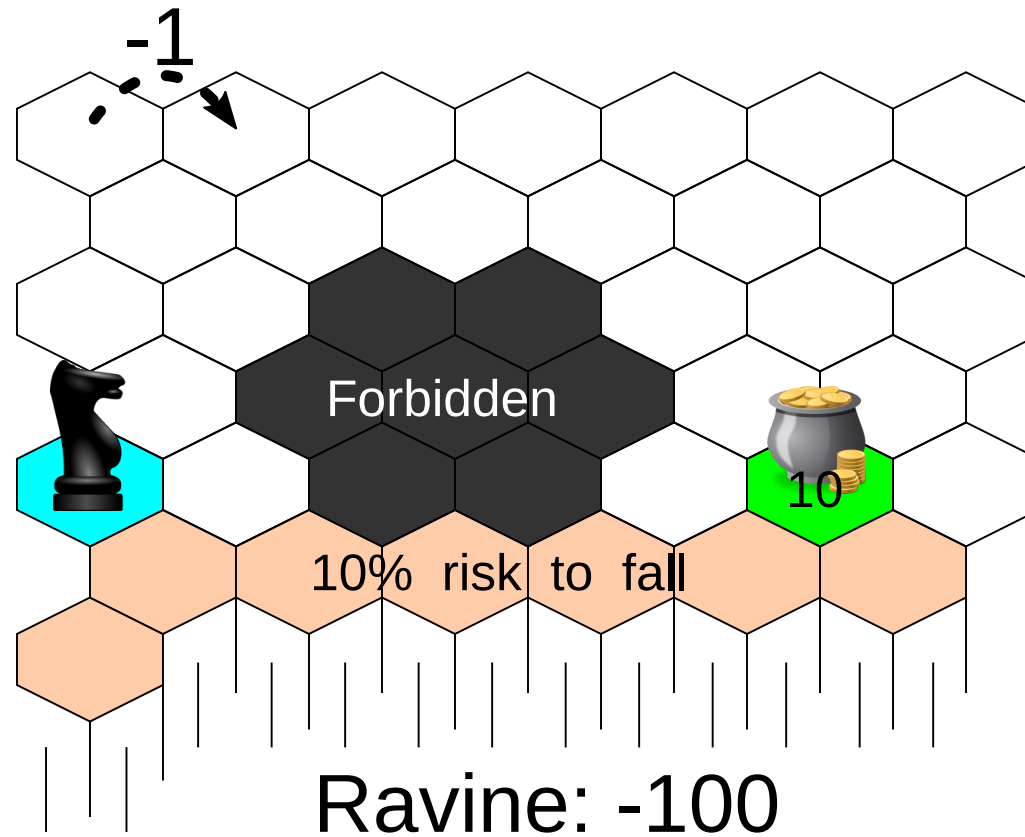
$$a^* = \arg \max_{a \in A} \left(R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') \times V(s') \right)$$

 ▶ Update $\pi(s)$ and $V(s)$ by considering action a^*

 ▶ Compute the delta value between the previous and the new $V(s)$

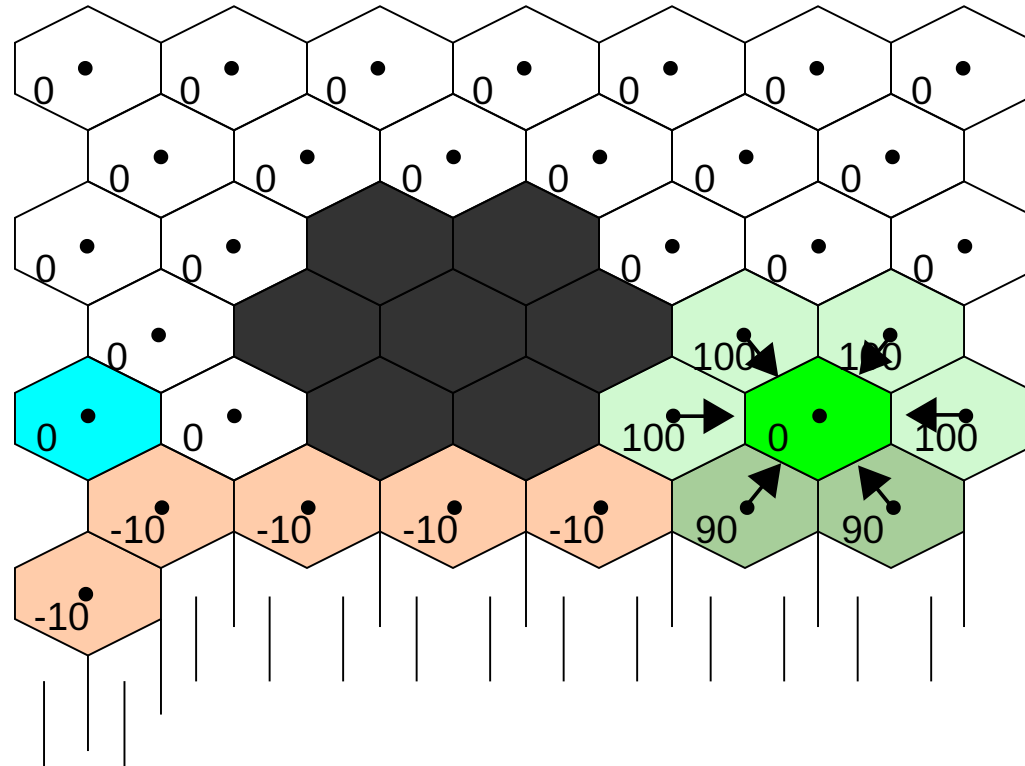
Output: an optimal π^* and associated V-values

Example: Dangerous move.




Probleme definition

Example: Dangerous move.



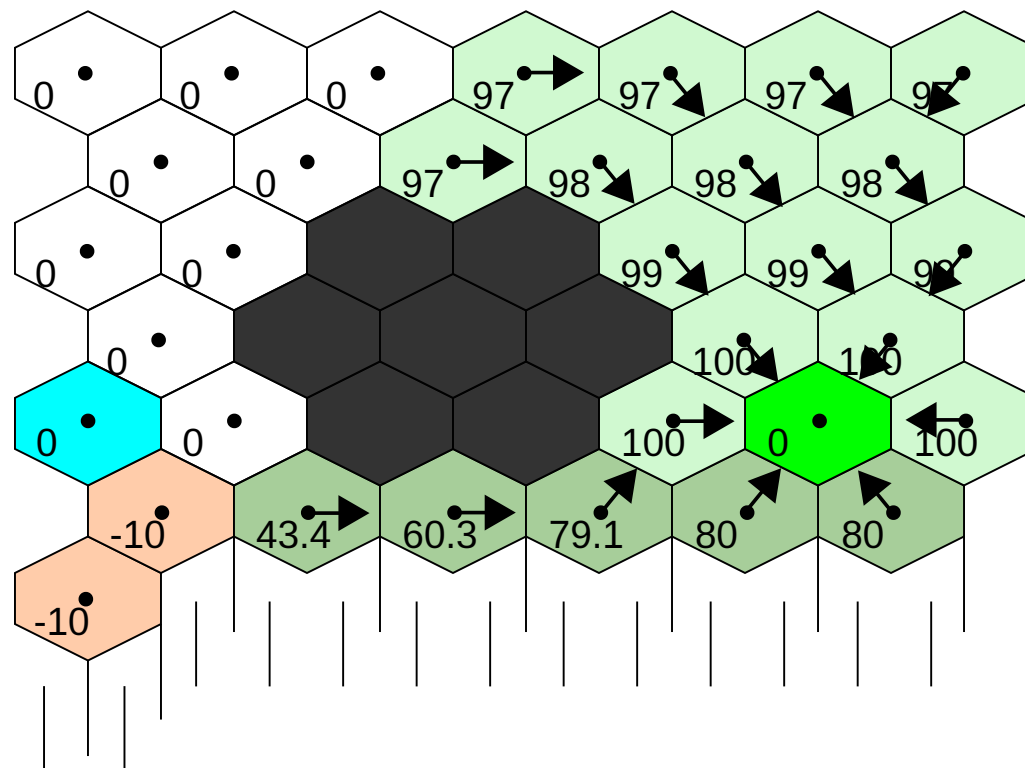
Value-Iteration: first iteration

Example: Dangerous move.




Value-Iteration: second iteration

Example: Dangerous move.




Value-Iteration: 4th iteration

Example: Dangerous move.




Value-Iteration: 6th iteration

Example: Dangerous move.



Value-Iteration: 7th iteration

Example: Dangerous move.



Value-Iteration: 8th iteration

Solving MDP: Policy Iteration

Input: an **MDP**: $\langle S, A, T, R \rangle$; precision error: ϵ ; discount factor: γ ; initial $V(s)$

1. Compute $\pi(s)$ according to $V(s)$, for each state $s \in S$
2. Repeat until $\pi(s)$ is stable:
 - Update $V(s)$ with $\pi(s)$ at ϵ error, for each state $s \in S$
 - Update $\pi(s)$ according to $V(s)$, for each state $s \in S$

Output: an optimal π^* and associated V-values

Ok now learn the model...

- ▶ Define the state-space (small but covering).
- ▶ Define the action-space.
- ▶ Explore the system:
 - Compute the average rewards $R(s, a)$.
 - Compute all transition probability $T(s, a, s')$

Learn the transition

The transition function is the core object to learn.

It is a 3-dimension structure of floating point values (probabilities).

$$|S|^2 \times |A| \text{ values.}$$

A simple game as **421** with **168** states and **8** actions
would requires **225 792** values.

Luky for us, in application, most of the transitions are null (ie. imposible).